


Chroma^[1]

Chen, Tianhao

2024/4/23

Illuminating protein space with a programmable generative model

[John B. Ingraham](#), [Max Baranov](#), [Zak Costello](#), [Karl W. Barber](#), [Wujie Wang](#), [Ahmed Ismail](#), [Vincent Frappier](#), [Dana M. Lord](#), [Christopher Ng-Thow-Hing](#), [Erik R. Van Vlack](#), [Shan Tie](#), [Vincent Xue](#), [Sarah C. Cowles](#), [Alan Leung](#), [João V. Rodrigues](#), [Claudio L. Morales-Perez](#), [Alex M. Ayoub](#), [Robin Green](#), [Katherine Puentes](#), [Frank Oplinger](#), [Nishant V. Panwar](#), [Fritz Obermeyer](#), [Adam R. Root](#), [Andrew L. Beam](#), ... [Gevorg Grigoryan](#) 

+ Show authors


Nature **623**, 1070–1078 (2023) | [Cite this article](#)

78k Accesses | **18** Citations | **297** Altmetric | [Metrics](#)

Contents


- Background (protein design)
- Chroma overview
- Diffusion (general)
- Design of Chroma's diffusion
- Conditional generation
- Wet lab experiments

Amino acids and proteins



Amino acid:

- **Ca**: alpha carbon
- NH₂: amino group
- R: sidechain
- COOH: carboxyl group




Formation of peptide bond

Peptide bond joins amino acids sequentially to form a protein sequence

Protein Modalities

MAKEDTLEFPGVVKELLPNAT
FRVELDNGHELIAVMAGKMRK
NRIRVLAGDKVQVEMTPYDLS
KGRINYRFK

Amino acid sequence



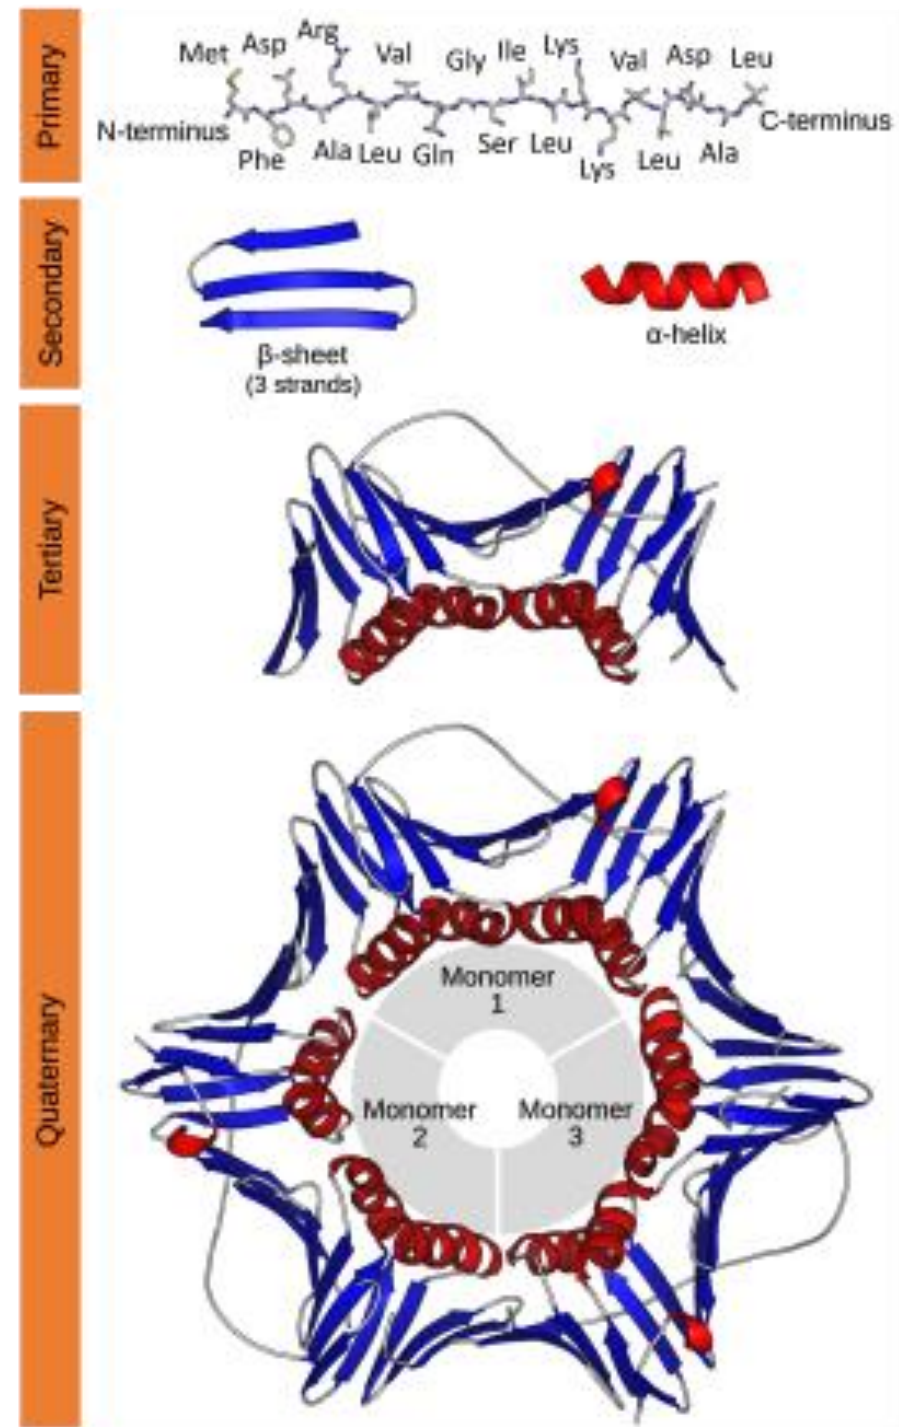
3D structure

Stabilizes the binding of IF-2
and IF-3 on the 30S subunit...
Belongs to the IF-1 family.
Subcellular location: Cytoplasm.

Biological functions and properties

Protein structure

- Primary
- Secondary
- Tertiary
- Quaternary




Protein design

- Folding
 - predict the structure that a given sequence might fold into
 - AlphaFold2^[2]
- Inverse folding
 - design sequence that folds in to given structure
 - ProteinMPNN^[3]
- **Structure / sequence design**
 - **Generate novel proteins** satisfying specific properties or constraints
 - Function, symmetry, substructure...

Chroma: overview


- Backbone generation
 - N-Ca-C=O
 - $\mathbf{x} \in \mathbb{R}^{4N \times 3}$
- Sequence and side-chain design
 - Based on generated backbone
 - Sequence: $\mathbf{s} \in [[20]]^N$
 - Amino acid type
 - Side-chain: $\chi \in (-\pi, \pi]^{4N}$
 - Torsion angles



$$\log p_{\theta}(\mathbf{x}, \mathbf{s}, \chi) = \underbrace{\log p_{\theta}(\mathbf{x})}_{\text{backbone likelihood}} + \underbrace{\log p_{\theta}(\mathbf{s}|\mathbf{x})}_{\text{sequence likelihood}} + \underbrace{\log p_{\theta}(\chi|\mathbf{x}, \mathbf{s})}_{\text{side-chain likelihood}}$$

diffusion Inverse folding

Diffusion^[4]



- Overview
 - Generate novel samples from a learned distribution
 - Learns a distribution by **learning to denoise** a noised data
 - Denoising process is broken down into many small steps for easier learning

- Forward process: gradually add (gaussian) noise to data

$$x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}z_t \longrightarrow x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}z$$

$$x_{t-1}|x_t, x_0 \sim N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$$

- Reverse process: gradually denoise noisy data

$$p_\theta(x_{t-1}|x_t) = N(x_{t-1}|\mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

$$p_\theta(x_T) = N(x_T|0, I) \quad \mu_\theta(x_t, t) = \tilde{\mu}_t(x_t, x_0);$$

- Sampling: start from pure noise, do reverse process

Diffusion: training

- Train on ELBO of $\log p_\theta(x_0)$
 - Regard x_1, \dots, x_T as latent variable
 - $p_\theta(x_{0:T})$ is the reverse process, parameterized by neural network
 - $q(x_{1:T}|x_0)$ is the forward process, replaces intractable $p_\theta(x_{1:T}|x_0)$

$$\mathbb{E}[-\log p_\theta(\mathbf{x}_0)] \leq \mathbb{E}_q \left[-\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right]$$

$$\mathbb{E}_q \left[\underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{L_T} + \sum_{t>1} \underbrace{D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))}_{L_{t-1}} - \underbrace{\log p_\theta(\mathbf{x}_0|\mathbf{x}_1)}_{L_0} \right]$$

Forward diffusion $dx_t = -\frac{1}{2}\beta(t)x_t dt + \sqrt{\beta(t)}dW_t$

Reverse diffusion $dy_t = \frac{1}{2}\beta(T-t)y_t dt + \beta(T-t) \underbrace{\nabla_{y_t} \ln \rho_{T-t}(y_t)}_{\text{score function}} dt + \sqrt{\beta(T-t)}dW_t$

Diffusion: training


- Train on ELBO of $\log p_\theta(x_0)$
 - Regard x_1, \dots, x_T as latent variable
 - $p_\theta(x_{0:T})$ is the reverse process, parameterized by neural network
 - $q(x_{1:T}|x_0)$ is the forward process, replaces intractable $p_\theta(x_{1:T}|x_0)$
- Essentially train on **denoising loss** / reconstruction loss

Algorithm 1 Training

```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
       $\nabla_\theta \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$   
6: until converged
```

Diffusion: why add noise

- **Intuition:** to learn a mapping from noise to data
- **Technically:** For better score estimation^{[15][16]}
 - score: $\nabla_x \log p(x)$
 - If $x \sim N(x_0, I)$, $\nabla_x \log p(x) \propto x_0 - x \rightarrow$ point towards high density region
- Diffusion is a **score-based** generative model
 - The score of noised $p_\theta(x_{t-1}|x_t)$ “point towards” x_{t-1}
$$p_\theta(x_{t-1}|x_t) = N(x_{t-1}|\mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$
 - Sampling of noiseless x_0 is achieved through a series of **score guidance**,
 $x_T \rightarrow x_{T-1} \rightarrow \dots \rightarrow x_0$
 - Tweedie’s formula: $\mathbb{E}(\mu|x) = x + \nabla_x \log p(x)$
 - $\mu \sim p(\mu)$ unobserved, $p(\mu)$ unknown
 - $x = \mu + \epsilon$, $\epsilon \sim N(0, I)$, observed



$$\nabla_x \log p(x) = \mathbb{E}(\mu|x) - x$$

point towards μ from x

Diffusion: why add noise

- **Score-based generative modeling**

- Score: $\nabla_x \log p(x)$

- How to sample with score: Langevin dynamics

$$\tilde{\mathbf{x}}_t = \tilde{\mathbf{x}}_{t-1} + \frac{\epsilon}{2} \nabla_{\mathbf{x}} \log p(\tilde{\mathbf{x}}_{t-1}) + \sqrt{\epsilon} \mathbf{z}_t$$

towards high likelihood avoid collapse at maxima

- How to estimate score: score matching

$$\frac{1}{2} \mathbb{E}_{p_{\text{data}}} [\|\mathbf{s}_{\boldsymbol{\theta}}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|_2^2]$$

Diffusion: why add noise

- **Challenges of noiseless score matching**


- Score undefined for most of the space
 - High dimensional data lie on low dimensional manifolds
 - Results difficulty in sampling
 - Noise perturbed score defined for entire space
- Score hard to learn for low density regions
 - Score matching accuracy depends on training data
 - Noise perturbed data is abundant

$$\frac{1}{2} \mathbb{E}_{p_{\text{data}}} [\| \mathbf{s}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x}) \|_2^2]$$

Low density region
→ little training data
→ poor estimation

Diffusion: why add noise

- **Challenges of noiseless score matching**
- Langevin sampling do not recover relative weights
 - Even with ground truth score
 - Annealed Langevin dynamics recovers relative weights



Diffusion: why add noise

- **Denoising score matching**^[17]

- Estimation of $\nabla_{\tilde{x}} \log q_{\sigma}(\tilde{x})$
 - $q_{\sigma}(\tilde{x}, x) = q_{\sigma}(\tilde{x}|x)q(x)$, $q_{\sigma}(\tilde{x}) = \int q_{\sigma}(\tilde{x}, x)dx$
 - Learn slightly perturbed score
- Explicit score matching ~ denoising score matching
 - Learn tractable score instead of intractable score
- Diffusion: DSM with multiple noise scales

$$J_{ESMq_{\sigma}}(\theta) = \mathbb{E}_{q_{\sigma}(\tilde{\mathbf{x}})} \left[\frac{1}{2} \left\| \psi(\tilde{\mathbf{x}}; \theta) - \frac{\partial \log q_{\sigma}(\tilde{\mathbf{x}})}{\partial \tilde{\mathbf{x}}} \right\|^2 \right]$$

ψ : parameterized score function

$$J_{DSMq_{\sigma}}(\theta) = \mathbb{E}_{q_{\sigma}(\mathbf{x}, \tilde{\mathbf{x}})} \left[\frac{1}{2} \left\| \psi(\tilde{\mathbf{x}}; \theta) - \frac{\partial \log q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})}{\partial \tilde{\mathbf{x}}} \right\|^2 \right]$$

derivation: exchange integral and partial

Diffusion: how to apply

- Choose variable / space to do diffusion on
 - Diffusion on raw data / latent features (latent diffusion^[5])
 - Diffusion on entire space / manifolds^[6] (e.g. on sphere)
 - Noise type: Gaussian, IGSO(3), Uniform...
- Forward schedule
 - Rate of noising data as time step increases
 - Linear schedule, Cosine schedule^[7], ...
 - Essentially weighs different SNR during training


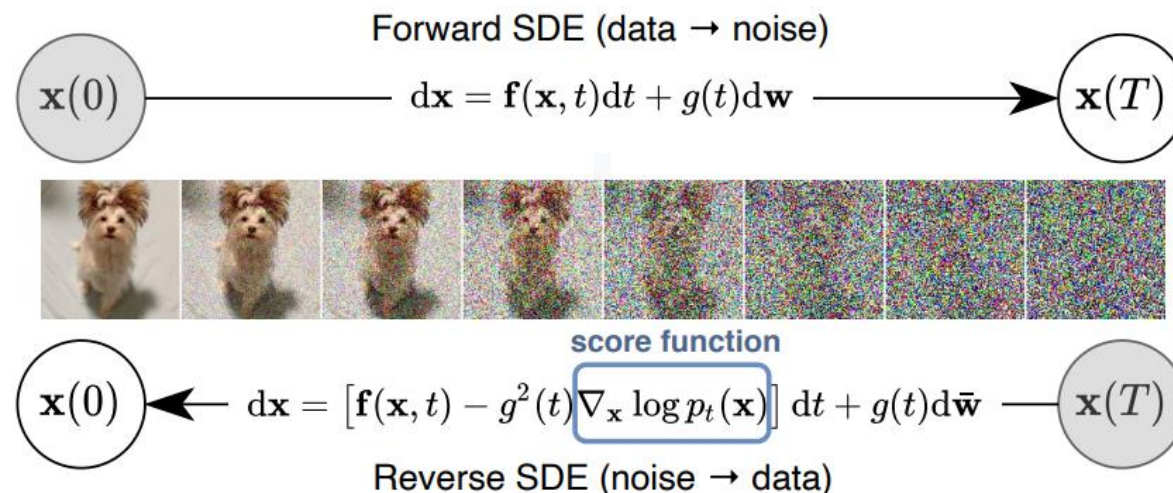


Figure 3. Latent samples from linear (top) and cosine (bottom)

Diffusion: how to apply

- Noise prediction network
 - A suitable neural network for diffusion variable
 - Input noisy data, output clean data / added noise
 - U-Net, DiT, GNN, ...
- Sampling procedure
 - Various SDE/ODE solvers based on learned **score function**^[8]




Diffusion: Chroma

- Diffusion on backbone coordinates
 - N-Ca-C=O
 - $\mathbf{x} \in \mathbb{R}^{4N \times 3}$
- Diffusion with correlated noise $p(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}; \alpha_t \mathbf{x}_0, \sigma_t^2 \mathbf{R} \mathbf{R}^\top)$
 - Intuition: removes simple correlation in data, and diffusion learns the residual
 - Correlation: backbone atoms of the same residue have similar coords
 - Analogy: RGB channels of an image are usually correlated



Amino acid diameter: 4~10 Å
Bond lengths: 1~2 Å



```
R:
tensor([ [ 1., 10., 0., 0.],
         [ 0., 10., 0., 0.],
         [ 0., 10., 1., 0.],
         [ 0., 10., 0., 1.] ])
```

```
(Pdb) z
tensor([ [-1.2461,  0.7056,  0.8065],
         [-0.2219, -0.5666, -1.0465],
         [ 0.0628,  0.4394, -0.0646],
         [ 0.8090, -0.4832,  0.8838] ])

(Pdb) R@z
tensor([ [-3.4653, -4.9601, -9.6582],
         [-2.2192, -5.6657, -10.4647],
         [-2.1564, -5.2262, -10.5294],
         [-1.4102, -6.1488, -9.5810] ])
```


Diffusion: Chroma

- Noise prediction network
 - Should be equivariant to **rotation and translation** (SE(3))
 - Rotations and translations should not affect denoise performance
- Equivariance
 - $f \circ g(x) = g \circ f(x)$, $\forall g \in \mathcal{G}$, then f is equivariant to \mathcal{G}
 - Invariance: $f \circ g(x) = f(x)$
- **SE(3)-equivariant GNN** as noise prediction network
 - $\text{denoise}(T @ X) = T @ \text{denoise}(X) \leftarrow x = a x_0 + b \epsilon$
 - T is an SE(3) transformation: $T(x) = Rx + t$


```
(Pdb) (T_X_denoised - TX_denoised).abs().mean().item()
0.08946086466312408
(Pdb) TX_denoised.abs().mean().item()
20.648839950561523
```

Equivariant GNN: How?

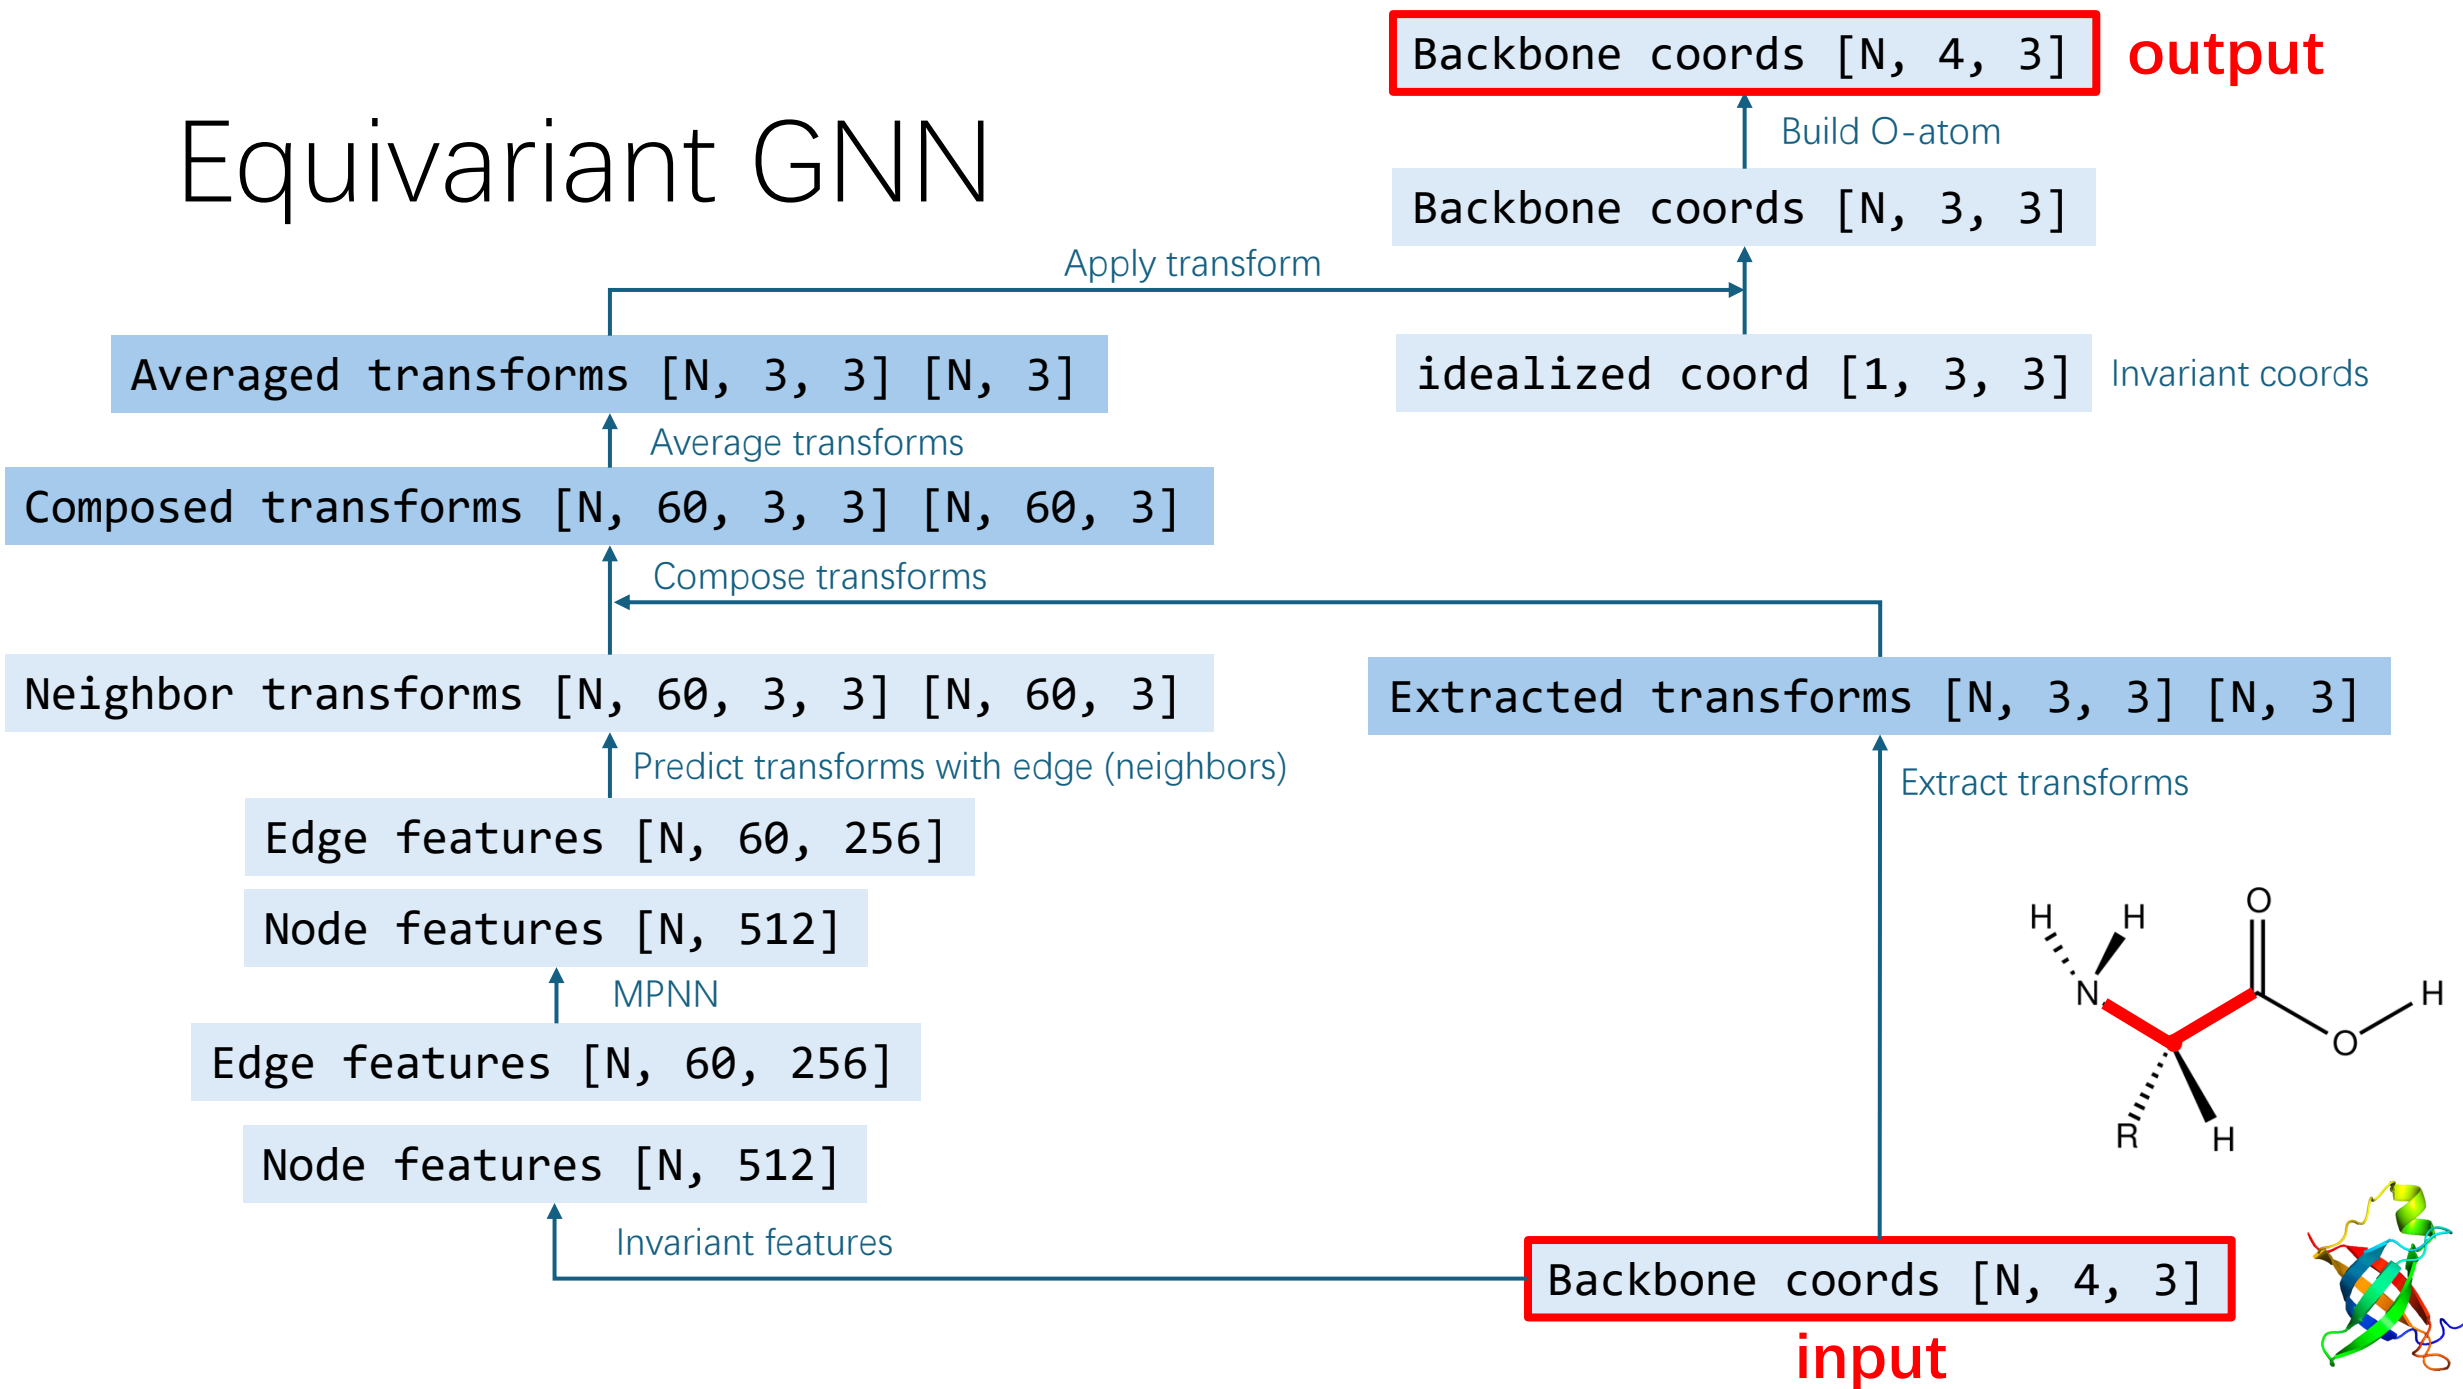
- Equivariant to SE(3) transformation on coordinates
- Method: **invariant features** + **equivariant operations**
- Initialize invariant node and edge features
 - Relative distances, internal angles^[9], bond lengths, protein diameters
 - Use MPNN to update initial features (concat neighbors, mlp, attn...)
- Update coords with equivariant operations
 - Extract, compose, average



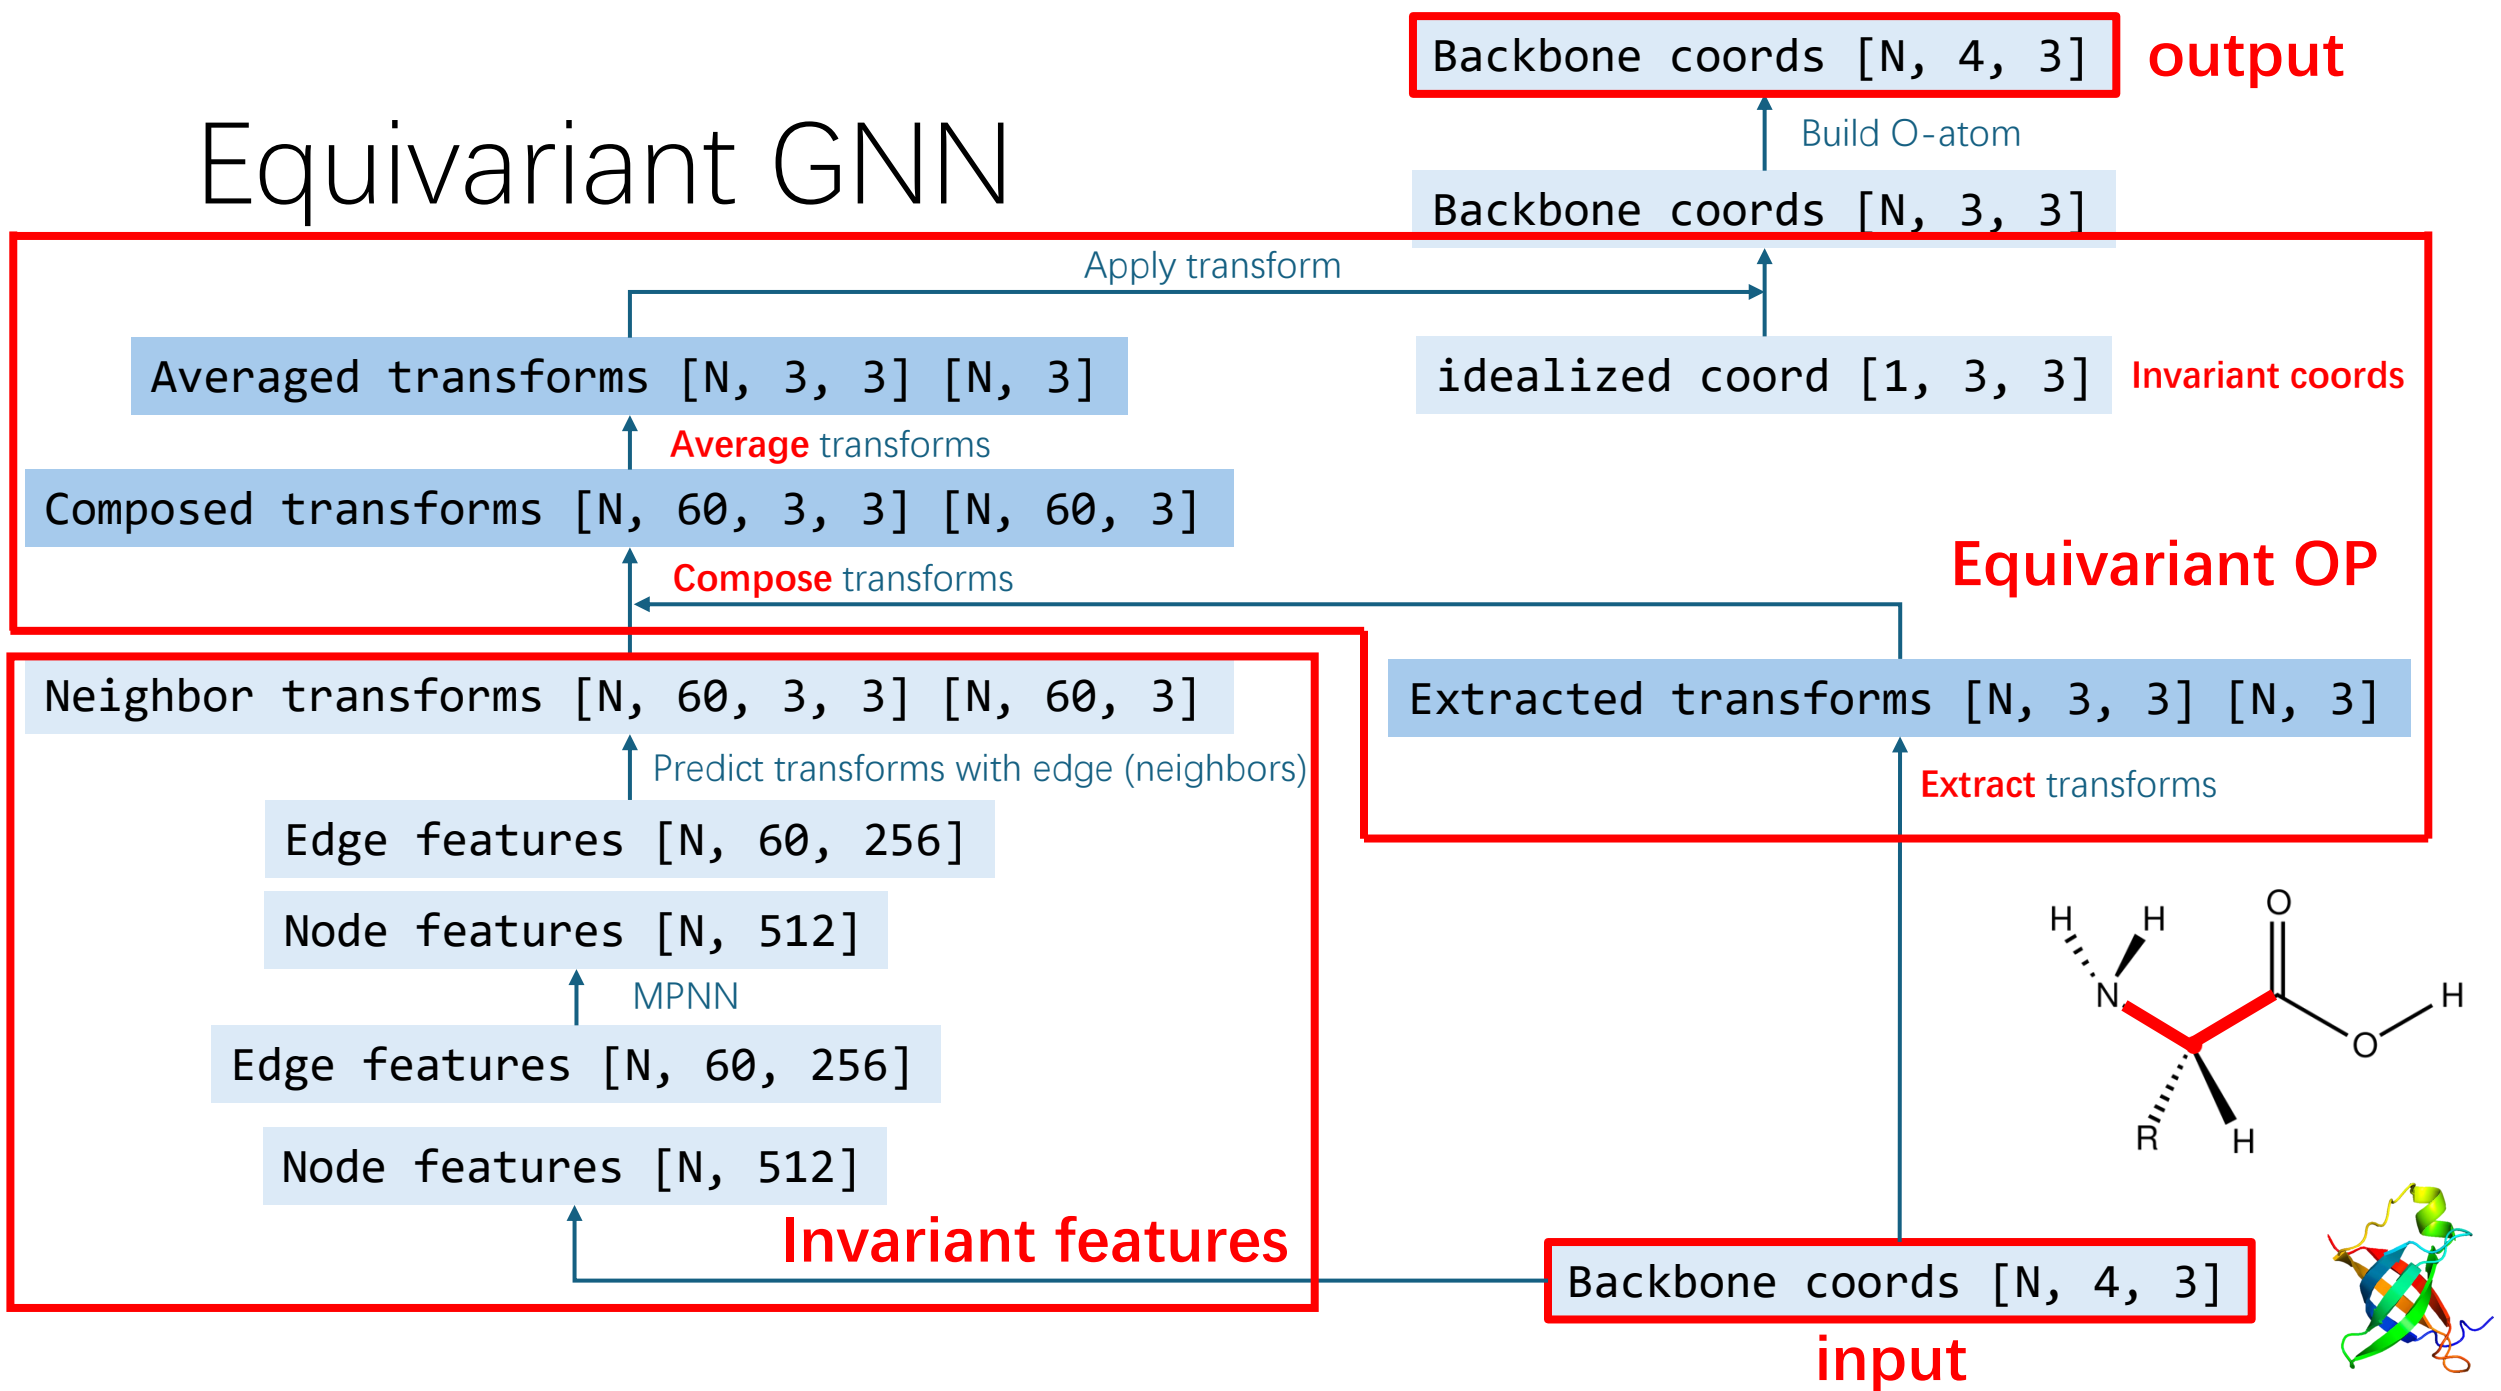
Angle	Description
ψ	Dihedral torsion about $N_i - C\alpha_i - C_i - N_{i+1}$
ω	Dihedral torsion about $C\alpha_i - C_i - N_{i+1} - C\alpha_{i+1}$
ϕ	Dihedral torsion about $C_i - N_{i+1} - C\alpha_{i+1} - C_{i+1}$
θ_1	Bond angle about $N_i - C\alpha_i - C_i$
θ_2	Bond angle about $C\alpha_i - C_i - N_{i+1}$
θ_3	Bond angle about $C_i - N_{i+1} - C\alpha_{i+1}$




Equivariant GNN



Equivariant GNN



GNN: Extract transforms



- **From backbone coords to rotations and translations**

- $\text{extract}(x) = (R, t) = T_{\text{extract}}$
- $[\text{N}, 4, 3] \rightarrow [\text{N}, 3, 3] + [\text{N}, 3]$
- $[\text{N}, 3, 3]$: 3x3 rotation matrix
 - $\text{Ca} \rightarrow \text{N}$ direction
 - $\text{Cross}(\text{Ca} \rightarrow \text{N}, \text{Ca} \rightarrow \text{C})$
 - Cross product of the above 2 directions
- $[\text{N}, 3]$: Ca coords

} Global rotation of N-Ca-C
} Global translation of N-Ca-C


- For reconstruction, use **canonical bond length and bond angles** (x^*)

- Diffusion noising process may produce variant bond lengths
- Only extract **direction** from noised coords, distance is ignored
- O atom coords are extrapolated from N-Ca-C coords

- $x \sim T_{\text{extract}}(x^*), \quad T(x) \sim T \circ T_{\text{extract}}(x^*)$


- **Equivariant:** $\text{extract}(T(x)) = T(\text{extract}(x))$

GNN: Compose transforms



- **Compose extracted transform and neighbor transforms**
- Extracted transform: (R_i, t_i)
 - $i = 1, \dots, \text{num_residue}$
 - Rotation and position of each individual amino acid
 - **Origin \rightarrow residue**
- Neighbor transforms: (R_{ji}, t_{ji})
 - $i = 1, \dots, \text{num_residue}, j = 1, \dots, \text{num_neighbor}$
 - **Relative** orientation and position of each amino acid from its neighbor's perspective
 - **Neighbor \rightarrow residue**
- Composed transforms: $(R_j R_{ji}, t_j + R_j t_{ji}) = (R_j, t_j) \circ (R_{ji}, t_{ji})$
 - **Absolute** orientation and position of each amino acid from its neighbor's perspective
 - **Origin \rightarrow neighbor \rightarrow residue (compose in **reverse** order)**

GNN: Compose transforms



- **Compose extracted transform and neighbor transforms**

- Intuition: from neighbors' relative view to global view

- $T_{\text{neighbor}} \rightarrow T_{\text{extract}} \circ T_{\text{neighbor}} \quad (T_{ji} \rightarrow T_j \circ T_{ji})$

- Note: T_{neighbor} is predicted from invariant edge features of x

- Previous slide: $\text{extract}(T(x)) = T(\text{extract}(x))$

- Compose in **reverse order** (for equivariance)

- $T_{\text{extract}} \circ T_{\text{neighbor}}$, for x . $(T_j \circ T_{ji})$

- $T \circ T_{\text{extract}} \circ T_{\text{neighbor}}$, for $T(x)$. $(T \circ T_j \circ T_{ji})$

Global transform of residue i
from neighbor j 's opinion

- Conclusion: Compose is **equivariant**

```
R_composed = R_a @ R_b
t_composed = t_a + (R_a @ t_b.unsqueeze(-1)).squeeze(-1)
return R_composed, t_composed
```

GNN: Average transforms

- **Average composed neighbor transforms (consensus structure)**

- $\text{avg}_R(\{R_i\})$ for $i = 1, \dots, \text{num_neighbor}$

- Weight $\{w_i\}$ from edge features
- Take SVD: $\sum w_i R_i = U D V^T$
- $\text{avg}_R(\{R_i\}) := U V^T \times \det(U V^T)$, (to ensure $\det = +1$)
- observation: $\text{avg}_R(\{R R_i\}) = R \text{avg}_R(\{R_i\})$

```
# Average rotation via SVD
R_avg_unc = (R * R_probs).sum(dim)
R_avg_unc = R_avg_unc + dither_eps * torch.randn_like(R_avg_unc)
U, S, Vh = torch.linalg.svd(R_avg_unc, full_matrices=True)
R_avg = U @ Vh
```

- $\text{avg}_t(\{t_i\})$

- Weight $\{w_i\}$ from edge features
- $\text{avg}_t(\{t_i\}) := \sum w_i t_i$
- observation: $\text{avg}_t(\{R t_i + t\}) = R \text{avg}_t(\{t_i\}) + t$

```
# Average translation.
t_avg = (t * t_probs).sum(dim)
```

- Conclusion: Average is **equivariant**

GNN: Average transforms

- **Average composed neighbor transforms (consensus structure)**
- Why SVD: consensus structure optimization
- $d(T, T_i) = \|R - R_i\|_F^2 + \|t - t_i\|^2$
- Find T minimize $\sum w_i d(T, T_i)$
- Average rotations:
 - $d(R, R_i) = \|R - R_i\|_F^2 = \text{tr}((R - R_i)^T (R - R_i))$
 - Minimize $\sum w_i d(R, R_i)$
 - \Leftrightarrow **Maximize $\text{tr}(R^T \bar{R})$, $\bar{R} = \sum w_i R_i$**

GNN: Average transforms

- **Average composed neighbor transforms (consensus structure)**
- Why SVD: consensus structure optimization
- **Maximize $\text{tr}(R^T \bar{R})$, $\bar{R} = \sum w_i R_i$**
- $\bar{R} = UDV^T = (UV^T)(VDV^T) = \tilde{U}P$
 - Polar decomposition: \tilde{U} orthogonal, $P \succcurlyeq 0$
- $R^T \bar{R} = (R^T \tilde{U})P = OP$, O orthogonal, $P \succcurlyeq 0$
- **Lemma:** $\text{tr}(OP) \leq \text{tr}(P)$, equal when $O|_P = I \iff R = \tilde{U} = UV^T$
 - Proof: using $P = VDV^T$, we have $\text{tr}(OP) = \sum d_i \langle Ov_i, v_i \rangle$


GNN: Average transforms

- **Average composed neighbor transforms (consensus structure)**
- Why SVD: consensus structure optimization

$$\mathbf{T}_i^* \leftarrow \arg \min_{\mathbf{T}_i} U(\{\mathbf{T}_i\}; \{w_{ij}, \mathbf{T}_{ij}\}).$$


$$\begin{aligned} U(\{\mathbf{T}_i\}; \{w_{ij}, \mathbf{T}_{ij}\}) &= \sum_{i,j} w_{ij} \|\mathbf{T}_i - \mathbf{T}_j \circ \mathbf{T}_{ji}\|^2 \\ &= \sum_{i,j} w_{ij} \|\mathbf{O}_i - \mathbf{O}_j \mathbf{O}_{ji}\|^2 + w_{ij} \|\mathbf{t}_i - (\mathbf{O}_j \mathbf{t}_{ji} + \mathbf{t}_j)\|^2. \end{aligned}$$

$$\mathbf{T}_i^* = \left(\text{Proj}_{\text{SO}(3)} \left(\sum_j p_{ij} \mathbf{O}_j \mathbf{O}_{ji} \right), \sum_j p_{ij} (\mathbf{O}_j \mathbf{t}_{ji} + \mathbf{t}_j) \right), \text{ where } p_{ij} = \frac{w_{ij}}{\sum_j w_{ij}}$$



GNN: all-together

- Input: x (noisy coords)
- Output: $T_{\text{GNN}}(x^*)$ (denoised coords)
 - x^* is idealized coords with Ca at origin, invariant to x
- $T_{\text{GNN}} = \underbrace{\text{average}}_{\text{equivariant}}(\underbrace{T_{\text{extract}}}_{\text{equivariant}} \circ \underbrace{T_{\text{neighbor}}}_{\text{equivariant}})$
 - Average, extract, compose are equivariant (equivariant OP)
 - T_{neighbor} is invariant (from invariant features)
- Conclusion: T_{GNN} is **equivariant**



```
(Pdb) self._t_atom  
tensor([[[[ 1.4590,  0.0000,  0.0000],  
            [ 0.0000,  0.0000,  0.0000],  
            [-0.5470,  0.0000, -1.4240]]]])
```

x^*

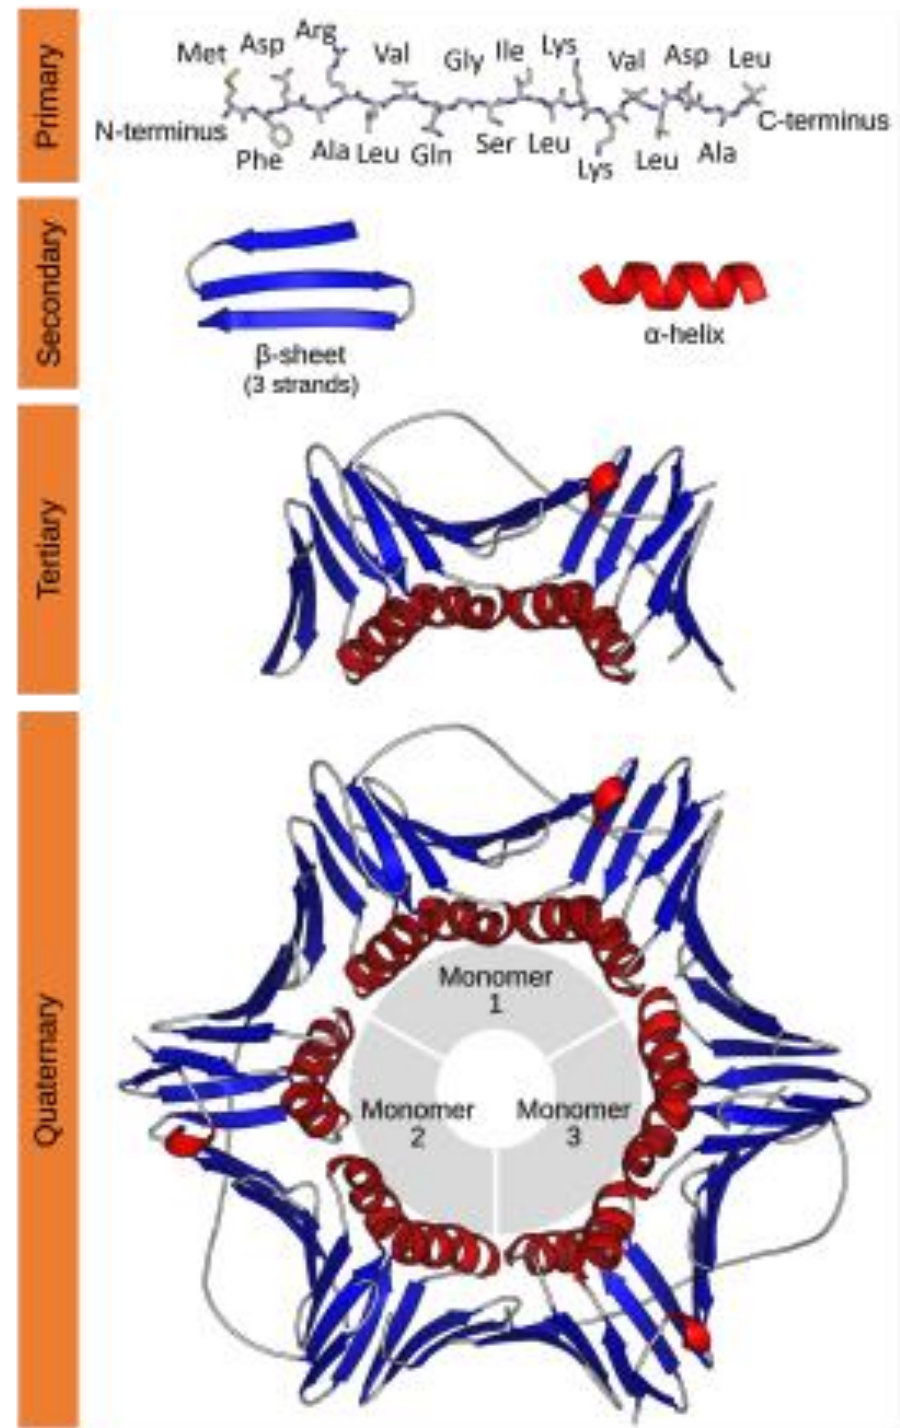
```
(Pdb) (T_X_denoised - TX_denoised).abs().mean().item()  
0.08946086466312408  
(Pdb) TX_denoised.abs().mean().item()  
20.648839950561523
```

Diffusion: Chroma training

- ELBO (backbone atom-wise MSE)
- Substructure MSE
 - Essentially weighted MSE emphasizing local substructure
 - Why: ELBO mainly focus on low-freq statistics
 - Analogy: SD generate fingers poorly
 - How: substructure regions of training data can be easily identified
- Relative distance matrix MSE

$$\mathcal{D}_{\text{distance}}(\mathbf{x}, \mathbf{x}') = \sum_{ij} (D_{ij}^{\text{CA}}(\mathbf{x}) - D_{ij}^{\text{CA}}(\mathbf{x}'))^2$$



(D: Internal distance)



Evaluation of unconditional samples

- **Refolding TM-score**
- Run inverse folding on generated backbone
- Predict structure of designed sequence with AlphaFold2
- Compare refolded structure and generated structure

d



Conditional generation

- General method
- Chroma's conditioners
 - Symmetry constraint
 - Substructure constraint
 - Protein class
 - Shape guidance
 - Text guidance

Conditional generation

- General method: **Conditional reverse process**
- Unconditional generation: sample from $p_\theta(x_{t-1}|x_t)$
- Conditional generation: sample from $p_\theta(x_{t-1}|x_t, y)$
 - y is the condition
- In reverse sampling, use **score** of sample distribution

$$\begin{aligned}\nabla_{\mathbf{x}} \log p_t(\mathbf{x}|\mathbf{y}) &= \nabla_{\mathbf{x}} \log \frac{p_t(\mathbf{x})p_t(\mathbf{y}|\mathbf{x})}{p_t(\mathbf{y})} \\ \text{Conditional score} \quad &= \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) + \nabla_{\mathbf{x}} \log p_t(\mathbf{y}|\mathbf{x}) - \cancel{\nabla_{\mathbf{x}} \log p_t(\mathbf{y})} \\ &= \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) + \boxed{\nabla_{\mathbf{x}} \log p_t(\mathbf{y}|\mathbf{x})}. \\ &\quad \text{Unconditional score} + \text{conditional gradient}\end{aligned}$$

Conditional generation

- General method: **Conditional reverse process**
- Unconditional generation: sample from $p_{\theta}(x_{t-1}|x_t)$
- Conditional generation: sample from $p_{\theta}(x_{t-1}|x_t, y)$
 - y is the condition
- In reverse sampling, use **score** of sample distribution
 - Combination of conditions

$$\nabla_{\mathbf{x}} \log p_t(\mathbf{x}|\mathbf{y}_1, \dots, \mathbf{y}_M) = \nabla_{\mathbf{x}} \log p_t(\mathbf{x}) + \sum_{i=1}^M \nabla_{\mathbf{x}} \log p_t(\mathbf{y}_i|\mathbf{x}).$$

Unconditional score + conditional gradients

Conditional generation

- General method: **Conditional reverse process**
- Unconditional generation: sample from $p_\theta(x_{t-1}|x_t)$
- Conditional generation: sample from $p_\theta(x_{t-1}|x_t, y)$
 - y is the condition
- Or use **projection** instead of $\nabla_x \log p_t(y|x)$
- Example^[10]: inpainting $\text{mask}(x_0^*)$

Noisy $x_t \rightarrow$ denoised x_0


$\rightarrow x_0 = \text{mask}(x_0^*) + (1 - \text{mask})(x_0)$

$\rightarrow x_{t-1} \sim q(x_{t-1}|x_t, x_0) \rightarrow$ denoised x_0

$\rightarrow \dots$

\leftarrow Unconditional update

\leftarrow Projection step


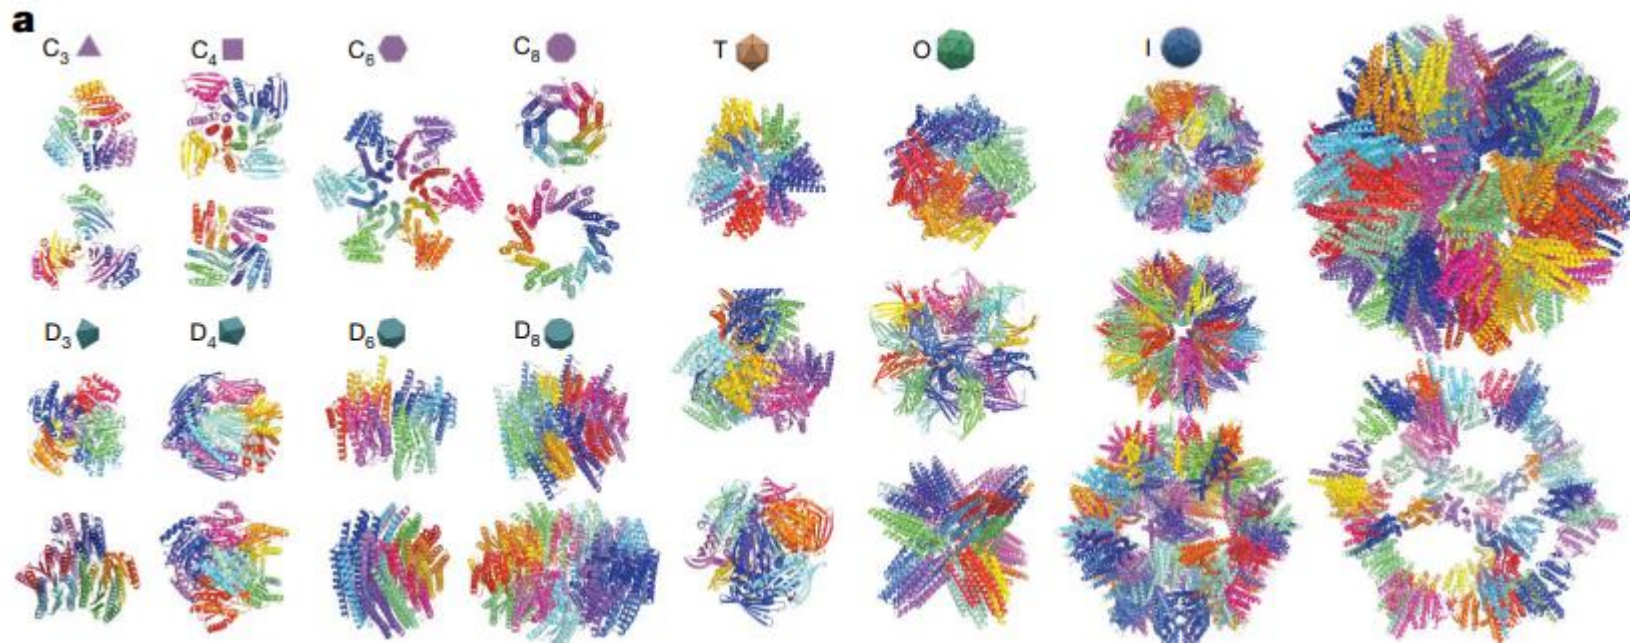


$\text{mask}(x_0^*)$

x_0^*

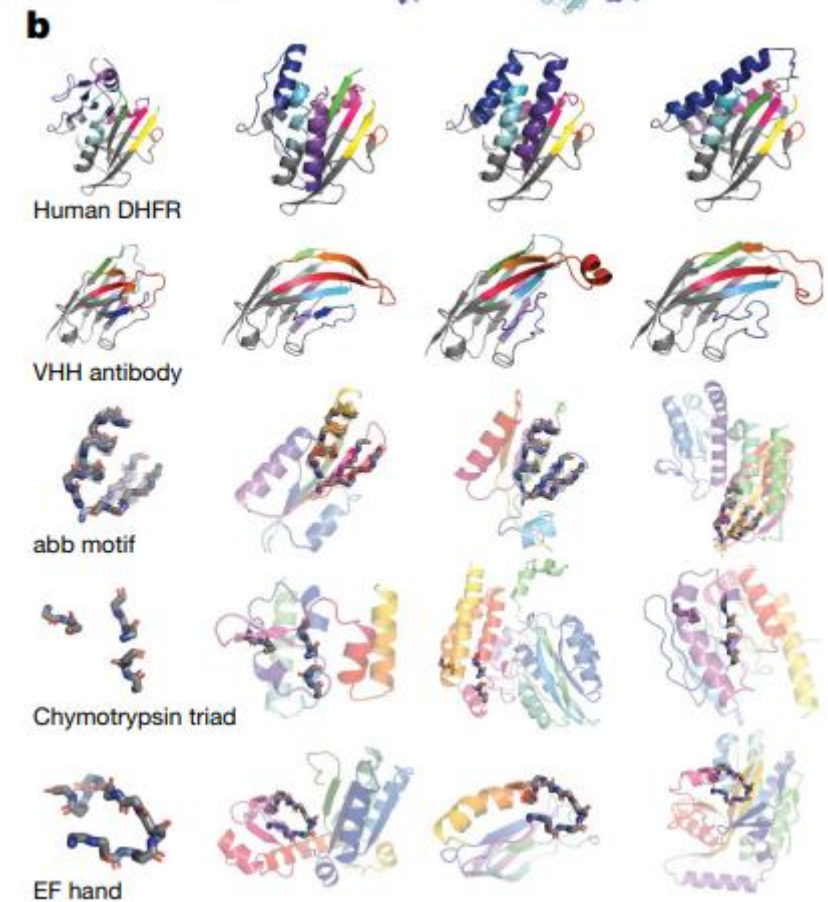

Conditional generation: Chroma

- **Symmetry constraint**
- Unconditional generation → projection
- Condition as rotation matrices




Conditional generation: Chroma

- **Substructure constraint** (inpainting)
- Unconditional generation → projection



Conditional generation: Chroma

- **Protein class condition** (classifier guidance^[11])
- $p(y|x_t)$: protein structure classifier
 - y : class label
 - x_t : noised backbone coords in the diffusion process
- Implementation of $p(y|x_t)$
 - GNN with classifier head

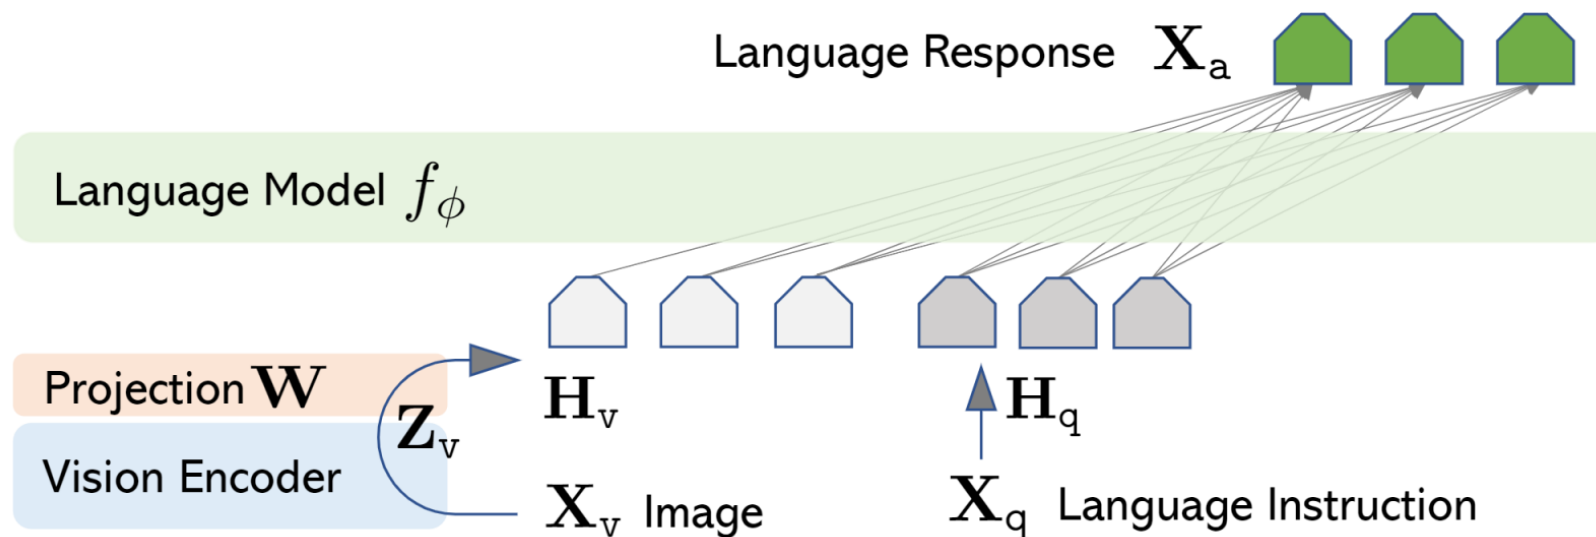


Conditional generation: Chroma

- **Text guidance**
- $p(y|x_t)$: backbone caption model
 - y : text description
 - x_t : noised backbone coords in the diffusion process
- Implementation of $p(y|x_t)$
 - A **multimodal language model**^[12]
 - Structure encoder: a GNN pretrained on classification tasks
 - Language model: GPT-Neo^[13], pretrained on arXiv, PubMed etc.
- Training data of $p(y|x_t)$: UniProt, PDB (protein caption data)
- Guidance mechanism
 - Use gradient $\nabla_{x_t} \log p(y|x_t)$ to guide denoising of x_t

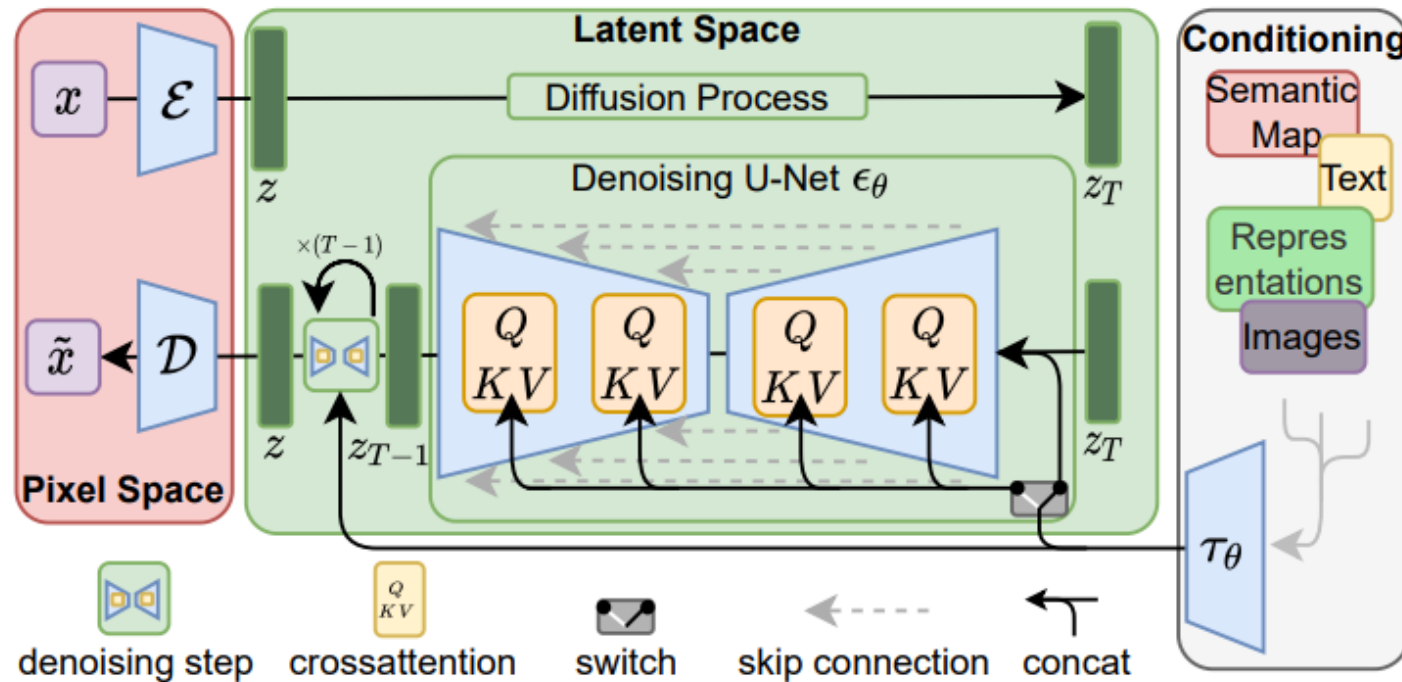
Conditional generation: Chroma

- **Text guidance**
- $p(y|x_t)$: backbone caption model
 - y : text description
 - x_t : noised backbone coords in the diffusion process
- Example of multimodal LLM: LLaVA^[12]



Conditional generation: Chroma

- **Text guidance** (different from SD)
- Stable diffusion text guidance
 - Uses cross-attention guidance, instead of external gradient




Conditional generation: Chroma

- **Shape guidance**
- Shape condition as point cloud
- $p(y|x_t)$: differentiable shape loss
 - N residues x_t , M points $y = r$
 - Wasserstein distance $K^W \in \mathbb{R}^{N \times M}$
 - Compare x_t and r inter-atomic distance
 - Gromov-Wasserstein distance $K^{GW} \in \mathbb{R}^{N \times M}$
 - Compare respective internal distances of x_t and r
 - More suitable for unaligned distribution in different space


$$\text{ShapeLoss}(\mathbf{x}, \mathbf{r}) = \sum_{i,j} \left(K_{ij}^{GW} + K_{ij}^W(\mathbf{x}, \mathbf{r}) \right) \|\mathbf{x}_i - \mathbf{r}_j\|,$$

inter-atomic distance



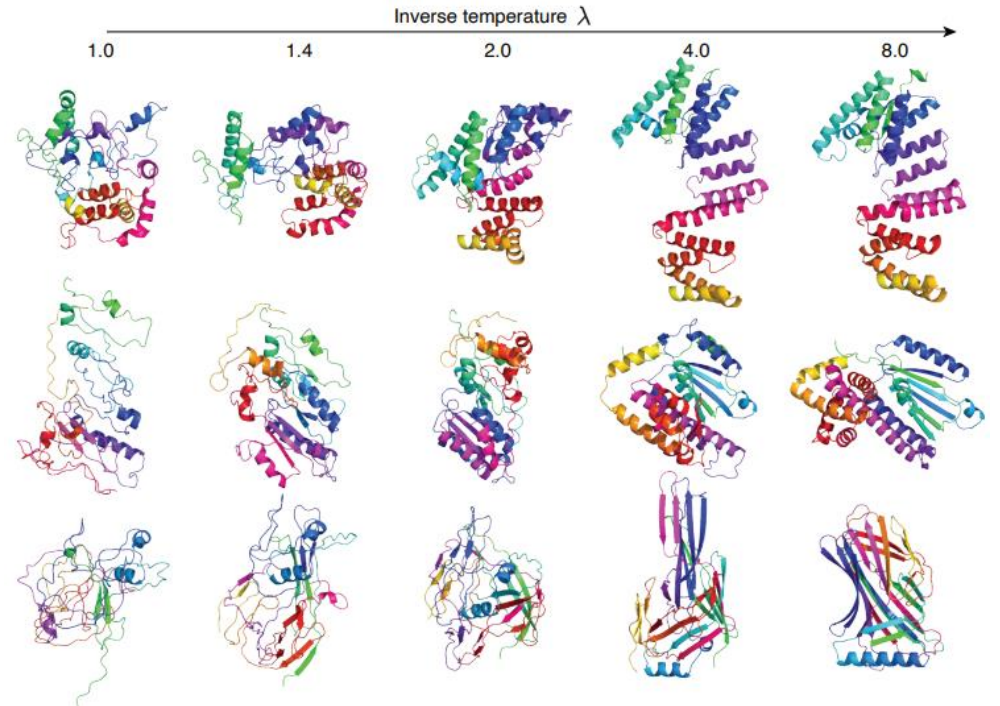
Low temperature sampling

- Goal: generate proteins with higher likelihood
- Motivation: higher likelihood usually means higher quality
- Temperature: trade-off between **quality and diversity**
 - higher temp \rightarrow higher variance, more diversified output
 - lower temp \rightarrow more concentrated
 - Analogy: greedy decoding, top-p sampling of LLM
- How
 - Scaling reverse process distribution $p_{\theta}(x_{t-1}|x_t)$
 - Classifier-free guidance^[14]

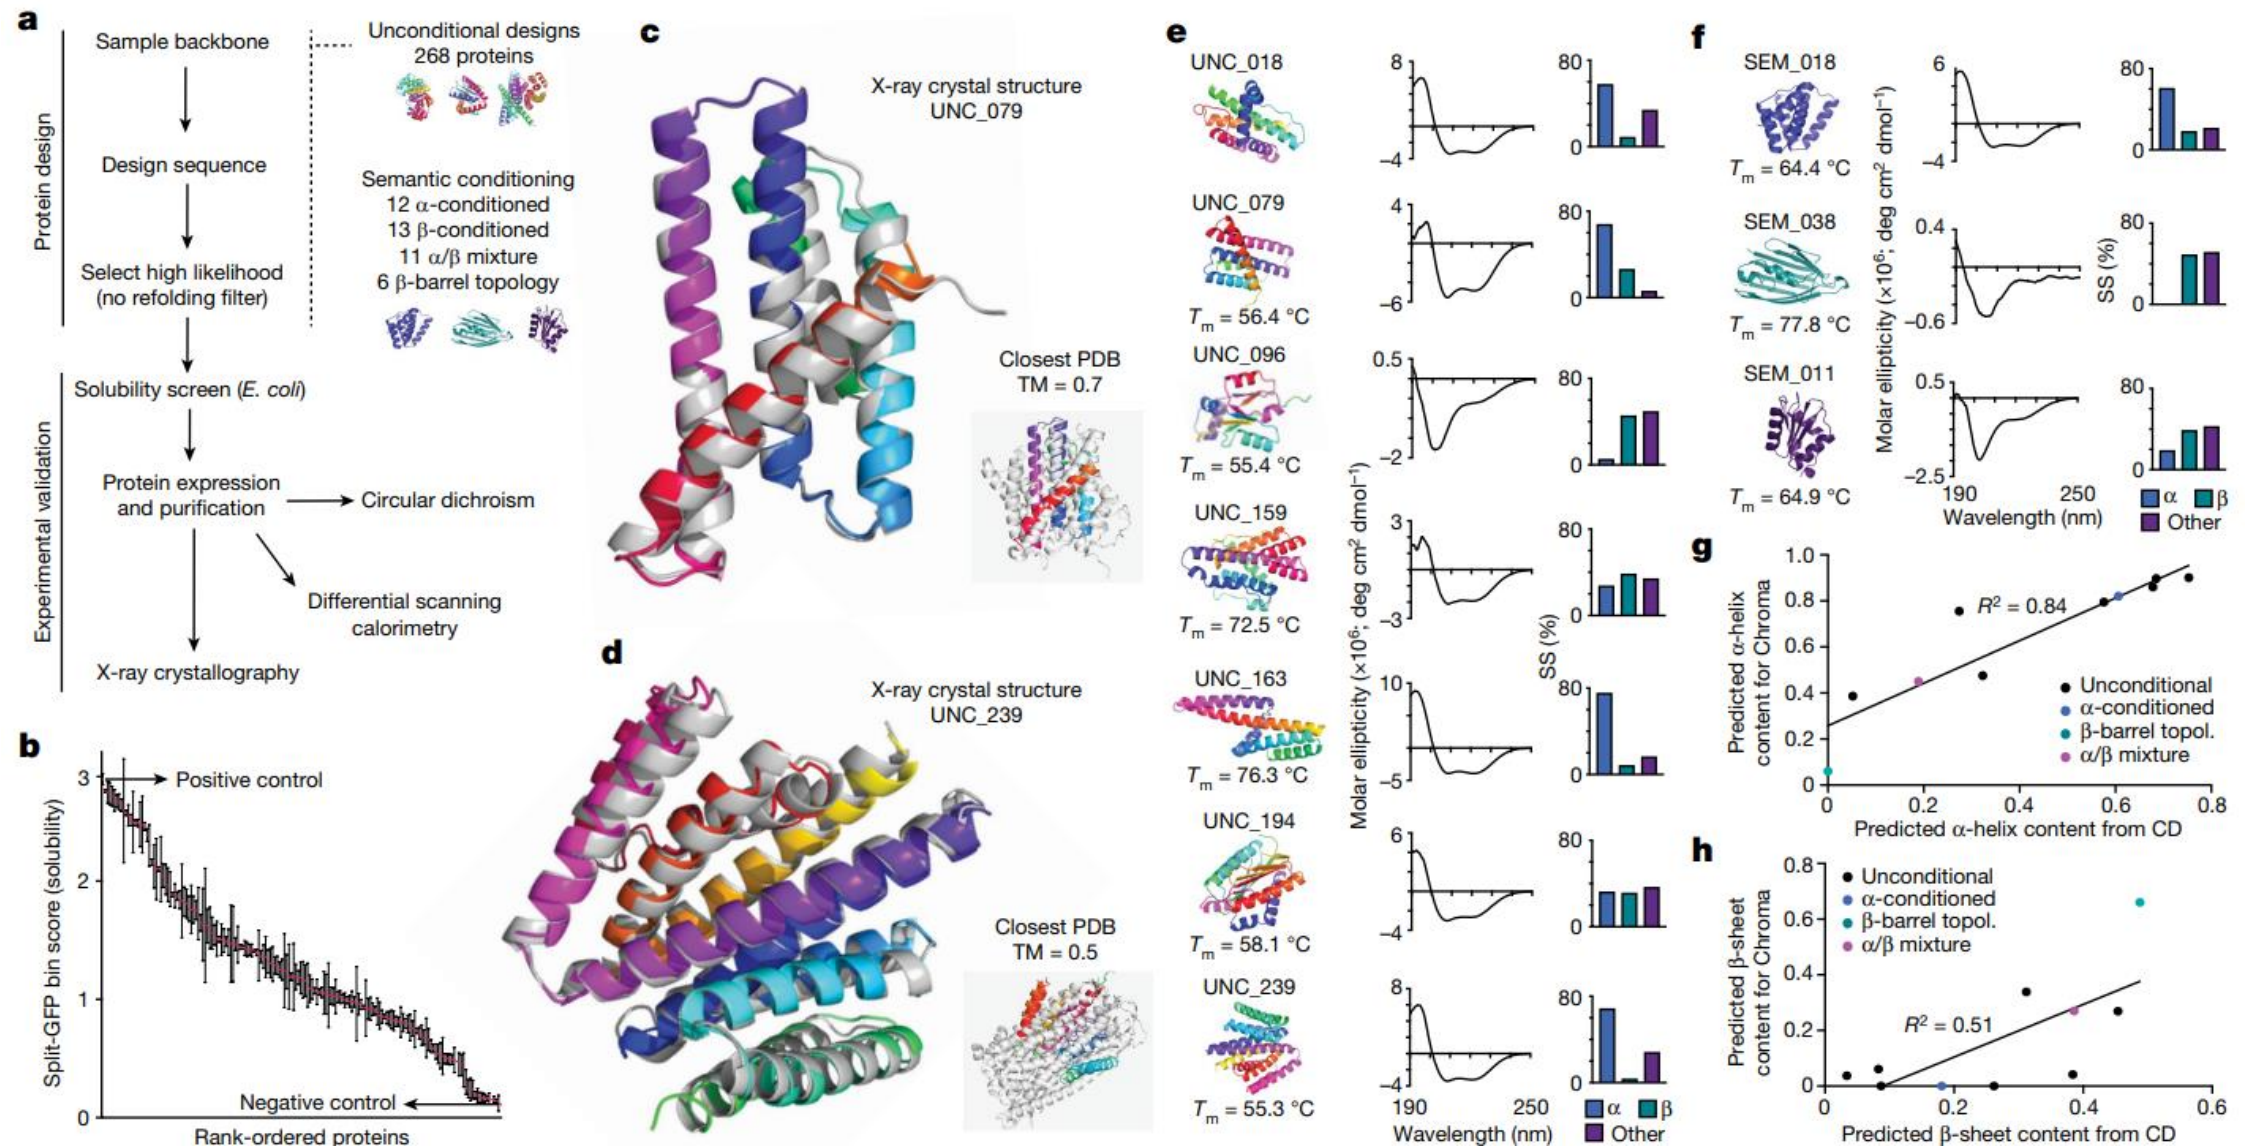


Low temperature sampling

- Scaling $p_{\theta}(x_{t-1}|x_t) \rightarrow \frac{1}{Z(\lambda)} p_{\theta}(x_{t-1}|x_t)^{\lambda}$
 - λ : inverse temperature
- Score is linearly scaled $\nabla_x \log p_{\theta}(x_{t-1}|x_t) \rightarrow \lambda \nabla_x \log p_{\theta}(x_{t-1}|x_t)$
- Scaling $p_{\theta}(x_0)$ is more accurate



Wet lab experiments



References

- [1] Ingraham, J.B., Baranov, M., Costello, Z. *et al.* Illuminating protein space with a programmable generative model. *Nature* **623**, 1070–1078 (2023).
- [2] Jumper, J., Evans, R., Pritzel, A. *et al.* Highly accurate protein structure prediction with AlphaFold. *Nature* **596**, 583–589 (2021).
- [3] J. Dauparas *et al.*, Robust deep learning–based protein sequence design using ProteinMPNN. *Science* **378**,49-56(2022).
- [4] Ho, J., Jain, A., Abbeel, P., Denoising Diffusion Probabilistic Models, *arXiv preprint*, 2006.11239
- [5] Robin, R., Andreas, B. Dominik, L. *et al.* High-Resolution Image Synthesis with Latent Diffusion Models, *arXiv preprint*, 2112.10752
- [6] Chin-Wei H.,Milad A., Avishek, B. *et al.* Riemannian Diffusion Models, *arXiv preprint*, 2208.07949
- [7] Alexander Quinn Nichol, Prafulla Dhariwal, Improved Denoising Diffusion Probabilistic Models, Proceedings of the 38th International Conference on Machine Learning, PMLR 139:8162-8171, 2021
- [8] Yang, S., Jascha S., Diederik K., Score-Based Generative Modeling through Stochastic Differential Equations}, International Conference on Learning Representations, 2021}
- [9] Wu, K.E., Yang, K.K., van den Berg, R. *et al.* Protein structure generation via folding diffusion. *Nat Commun* **15**, 1059 (2024).
- [10] Hyungjin, C., Jeongsol, K., Michael, T.M. *et al.* Diffusion Posterior Sampling for General Noisy Inverse Problems}, The Eleventh International Conference on Learning Representations, 2023
- [11] Prafulla Dhariwal and Alex Nichol. Diffusion models beat GANs on image synthesis. *arXiv preprint*, 2105.05233
- [12] Liu, H., Li, C., Wu, Q. and Lee, Y., Visual Instruction Tuning, NeurIPS, 2023
- [13] Black, S., Leo, G., Wang, P., Leahy, C., & Biderman, S. (2021). GPT-Neo: Large Scale Autoregressive Language Modeling with Mesh-Tensorflow (1.0). Zenodo.
- [14] Ho, J., Tim Salimans Classifier-Free Diffusion Guidance, *arXiv preprint*, 2207.12598
- [15] Yang Song, and Stefano Ermon. Generative Modeling by Estimating Gradients of the Data Distribution. NeurIPS 2019
- [16] P. Vincent. A connection between score matching and denoising autoencoders. *Neural computation*, 23(7):1661–1674, 2011.