Linear Regression and Classification

Yuan Yao

Department of Mathematics
Hong Kong University of Science and Technology

Most of the materials here are from Chapter 3-4 of Introduction to Statistical learning by Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani.

Spring, 2022

Linear Regression
    The least squares estimation
    The statistical properties of the least squares estimates.

Linear Classification
    Logistic Regression.
    Linear Discriminant Analysis
    The ROC curve

Outline

Linear Regression
The least squares estimation
The statistical properties of the least squares estimates.

Linear Classification
Logistic Regression.
Linear Discriminant Analysis
The ROC curve

Example: Advertising data

The data contains 200 observations.
Sample size: $n = 200$.
Sales: y_i, $i = 1..., n$.
TV (bugdets): $x_{i1}, i = 1, ..., n$.
Radio (budgets): $x_{i2}, i = 1, ..., n$.
Newspaper (budgets): $x_{i3}, i = 1, ..., n$.
Dimensionality: $p = 3$.

Example: Advertising data

Linear models formulation

► Consider linear regression model:

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i, \qquad i = 1, ..., n \quad (2.1)$$

where $y_i, x_i = (x_{i1}, ..., x_{ip})$ are the i-th observation of the response and covariates.

► Responses are sometimes called dependent variables or outputs;

► Covariates called independent variables, or predictors or features or inputs or regressors.

► Noise ϵ_i is independent, of zero mean, and fixed but unknown variance, e.g. Gaussian noise $\mathcal{N}(0, \sigma^2)$

Example: Advertising data

Now, we consider three covariates: TV, radio and newspapers.
The number of covariates (predictors, or features): $p = 3$.
The linear regression model is

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \epsilon_i, \qquad i = 1, ..., n$$

Estimating the coefficient by the least squares

Minimizing the sum of squares of error (Gauss'1795) [essentially as $-\log$ **likelihood** of normal distribution]:

$$\min_{\beta_0,\ldots,\beta_3} \sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_{i1} - \beta_2 x_{i2} - \beta_3 x_{i3})^2.$$

Figure: Carl Friedrich Gauss

Notations

With slight abuse of notation, in this chapter, we use

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{12} & ... & x_{1p} \\ 1 & x_{21} & x_{22} & ... & x_{2p} \\ \vdots & \vdots & \vdots & ... & \vdots \\ 1 & x_{n1} & x_{n2} & ... & x_{np} \end{pmatrix}$$

$$= \begin{pmatrix} \mathbf{1} \vdots \mathbf{x}_1 \vdots \mathbf{x}_2 \vdots \cdots \vdots \mathbf{x}_p \end{pmatrix}.$$

Here a column of ones, $\mathbf{1}$, is added, which corresponds to the intercept β_0. Then $\mathbf{X}$ is a n by $p+1$ matrix.

Recall that

$$\mathbf{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} \quad \beta = \begin{pmatrix} \beta_0 \\ \vdots \\ \beta_p \end{pmatrix}, \quad \epsilon = \begin{pmatrix} \epsilon_1 \\ \vdots \\ \epsilon_n \end{pmatrix}, \quad \mathbf{x}_j = \begin{pmatrix} x_{1j} \\ \vdots \\ x_{nj} \end{pmatrix}$$

The least squares criterion

The least squares criterion is try to minimize the sum of squares:

$$\sum_{i=1}^{n}(y_i - \beta_0 - \beta_1 x_{i1} - \cdots - \beta_p x_{ip})^2.$$

Using matrix algebra, the above sum of squares is

$$\|\mathbf{y} - \mathbf{X}\beta\|^2 = (\mathbf{y} - \mathbf{X}\beta)^T(\mathbf{y} - \mathbf{X}\beta).$$

The LSE, fitted values and residuals

By some linear algebra calcuation, the least squares estimator of β is then

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

Then

$$\hat{\mathbf{y}} = \mathbf{X}\hat{\beta}$$

is called the fitted values; viewed as the predicted values of the responses based on the linear model.

$$\hat{\epsilon} = \mathbf{y} - \hat{\mathbf{y}}$$

are called residuals. The sum of squares of these residuals

$$RSS = \sum_{i=1}^{n} \hat{\epsilon}_i^2 = \sum_{i=1}^{n} (y_i - \hat{y}_i)^2 = \|\mathbf{y} - \hat{\mathbf{y}}\|^2.$$

Orthogonality

▶ The residual $\hat{\epsilon}$ is orthogonal to all columns of $\mathbf{X}$, i.e, all $\mathbf{1}, \mathbf{x}_1, ..., \mathbf{x}_p$. This can be seen by

$$\mathbf{X}^T \hat{\epsilon} = \mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \hat{\beta}$$
$$= \mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{H} \mathbf{y} = 0.$$

where $\mathbf{H} = \mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T = \mathbf{H}^2$ is the projection.

▶ The residual vector $\hat{\epsilon}$ is orthogonal to the hyperplane formed by vectors $\mathbf{1}, \mathbf{x}_1, ..., \mathbf{x}_p$ in n dimensional real space.

The least squares projection

FIGURE 3.2. *The N-dimensional geometry of least squares regression with two predictors. The outcome vector* **y** *is orthogonally projected onto the hyperplane spanned by the input vectors* $\mathbf{x}_1$ *and* $\mathbf{x}_2$*. The projection* $\hat{\mathbf{y}}$ *represents the vector of the least squares predictions*

Gauss-Markov Theorem

Theorem (Gauss-Markov Theorem). Assume the noise ϵ_i is of
- mean zero, i.e. $\mathbb{E}(\epsilon_i) = 0$,
- homoscedastic in variance, i.e. $\mathbb{E}(\epsilon_i^2) = \sigma^2 < \infty)$,
- uncorrelated, i.e. $\mathbb{E}(\epsilon_i \epsilon_j) = 0)$.

Then the least square estimate is the *Best Linear Unbiased Estimate (BLUE)*, i.e. among all linear unbiased estimates, the least squares estimate has the smallest variance, thus smallest mean squared error.

- Note: no Gaussian distribution assumption on noise!

Figure: Left: Carl Friedrich Gauss; Right: Andrey Andreyevich Markov

Proof

- Let $\mathbf{A} = (\mathbf{a}_1, \mathbf{a}_2, ..., \mathbf{a}_n)$. Then, unbiased linear estimate of β is

$$\mathbf{Ay}$$

  with
    - mean $\mathbb{E}[\mathbf{Ay}] = \mathbf{AX}\beta = \beta$ by the unbiasedness,
    - and variance matrix $\mathbf{AA}^T$.

- Write $\mathbf{A} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T + \mathbf{D}$. Then, $\mathbf{DX} = 0$.

- Then,
$$\mathbf{AA}^T = (\mathbf{X}^T\mathbf{X})^{-1} + \mathbf{DD}^T \succeq (\mathbf{X}^T\mathbf{X})^{-1}.$$

  Here the inequality is for symmetric matrices, i.e., $\mathbf{A} \succeq \mathbf{B}$ is defined as $\mathbf{A} - \mathbf{B}$ is nonnegative definite. $\qquad\square$

Result of the estimation

TABLE 3.9. (from ISLR) The advertising data: coefficients of the LSE for the regression on number of units sold on TV, radio and newspaper advertising budgets.

	Coefficient	Std.error	t-statistic	p-value
Intercept	2.939	0.3119	9.42	< 0.0001
TV	0.046	0.0014	32.81	< 0.0001
radio	0.189	0.0086	21.89	< 0.0001
newspaper	-0.001	0.0059	-0.18	0.8599

Statistical properties of LSE

Assume $\epsilon_i \sim N(0, \sigma^2)$,

$$\hat{\beta} \sim N(\beta, \sigma^2(\mathbf{X}^T\mathbf{X})^{-1});$$

$$\text{RSS} = \sum_{i=1}^{n} \hat{\epsilon}_i^2 = \sum_{i=1}^{n}(y_i - \hat{y}_i)^2 \sim \sigma^2 \chi^2_{n-p-1}$$

$$\hat{\beta} \text{ and RSS are independent}$$

$$s^2 = \text{RSS}/(n-p-1) \text{ unbiased estimate of } \sigma^2$$

$$\frac{\hat{\beta}_j - \beta_j}{s\sqrt{c_{jj}}} \sim t_{n-p-1}$$

$$\frac{(\hat{\beta}-\beta)^T(\mathbf{X}^T\mathbf{X})(\hat{\beta}-\beta)/p}{s^2} \sim F_{p+1,n-p-1}$$

where $c_{00}, c_{11}, ..., c_{pp}$ are the diagonal elements of $(\mathbf{X}^T\mathbf{X})^{-1}$.

Confidence intervals

For example,

$$\hat{\beta}_j \pm t_{n-p-1}(\alpha/2)s\sqrt{c_{jj}}$$

is a confidence interval for β_j at confidence level $1 - \alpha$. Here $t_{n-p-1}(\alpha/2)$ is the $1 - \alpha/2$ percentile of the t-distribution with degree of freedom $n - p - 1$.

- the j-th p-value tells the probability of seeing $\hat{\beta}_j$ with assumption $\beta_j = 0$, useful for variable selection or elimination

Prediction interval

For a given value of input $\mathbf{x}$ which is a $p+1$ vector (the first component is constant 1), its mean response is $\beta^T \mathbf{x}$. The confidence interval for this mean response, also called prediction interval, is

$$\hat{\beta}^T \mathbf{x} \pm t_{n-p-1}(\alpha/2)s\sqrt{\mathbf{x}^T(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{x}}$$

- The confidence interval for β_j is a special case of the above formula by taking $\mathbf{x}$ as a vector that all zero except the $(j+1)$ entry corresponding β_j. (Because of β_0, β_j is at the $j+1$th position of $\hat{\beta}$.)

Variable selection

- We may be concerned with a subset of the p variables are irrelevant with the response.

- Let the subset be denoted as $A = \{i_1, ..., i_r\}$, where $r \leq p$. Then, the null hypothesis is

$$H_0 : \beta_{i_1} = \beta_{i_2} = \cdots = \beta_{i_r} = 0,$$

which again is equivalent to

$$H_0 : \mathbf{P}(\mathbf{y}) \in \mathcal{L}(A^c),$$

where $\mathcal{L}(A)$ is the linear space in R^n spanned by $\mathbf{x}_{i_1}, ..., \mathbf{x}_{i_r}$, which is of r dimension.

The F-test

$$F = \frac{\|\hat{\mathbf{y}} - \hat{\mathbf{y}}_0\|^2/(p+1-r)}{\|\mathbf{y} - \hat{\mathbf{y}}\|^2/(n-p-1)} \sim F_{p+1-r, n-p-r}$$

This F-statistic is used to test the hypothesis that
$H_0 : \mathbf{P}(\mathbf{y}) \in \mathcal{L}(A^c)$, against the alternative $H_a :$ otherwise.

- The smaller $Prob(F)$, the more significant of H_a
- The commonly considered hypothesis, $H_0 : \beta_1 = \cdots = \beta_p = 0$ can be formulated as $H_0 : \mathbf{P}(\mathbf{y}) \in \mathcal{L}(\mathbf{1})$, where $\mathcal{L}(\mathbf{1})$ represent the linear space of a single vector $\mathbf{1}$.

Outline

Examples: The default data

- Simulated data: 10000 individuals.
- Three predictors: income, balance (monthly), and student (Yes or No)
- One output: Default (Yes or No).
- Judge if one fails to pay the credit card debt (default).

Figure: FIGURE 4.1. The Default data set. Left: The annual incomes and monthly credit card balances of a number of individuals. The individuals who defaulted on their credit card payments are shown in orange, and those who did not are shown in blue. Strong relation between balance and default while a weaker relation between income and default.

The special case of binary response

- Consider the output is binary: two class,
- Code the response into 0 and 1 and apply linear regression produce the same result as linear discriminant anlaysis.
- Not the case for output with more than two classes.

Example: The default data

Figure: Left: linear regresion; Right: logistic regression

The setup for binary output

- The training data: $(\mathbf{x}_i, y_i)$: $i = 1, ..., n$.
- $y_i = 1$ for class 1 and $y_i = 0$ for class 0.
- $\mathbf{x}_i = (1, x_{i1}, ..., x_{ip})$ are $p + 1$ vectors with actually p inputs.
- If instead consider linear regression model is

$$y_i = \beta^T \mathbf{x}_i + \epsilon_i$$

  β can be estimated by the least squares, and $\hat{\beta}^T \mathbf{x}_i$ is the predictor of y_i.
- Key idea: should focus on predicting the probability of the classes.
- Using $P(y = 1|\mathbf{x}) = \beta^T \mathbf{x}$ is not appropriate.

The logistic regression model

▶ Assume

$$P(y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\beta^T \mathbf{x})}$$

As a result,

$$P(y = 0|\mathbf{x}) = 1 - P(y = 1|\mathbf{x}) = \frac{1}{1 + \exp(\beta^T \mathbf{x})}$$

▶

$$\log(\frac{P(y = 1|\mathbf{x})}{P(y = 0|\mathbf{x})}) = \beta^T \mathbf{x}$$

This is called log-odds or logit. And

$$\frac{P(y = 1|\mathbf{x})}{P(y = 0|\mathbf{x})}$$

is called odds.

▶ Interpretation: one unit increase in variable x_j, increases the log-odds of class 1 by β_j.

The maximum likelihood estimation

▶ Recall that, the likelihood is the joint probability function of joint density function of the data.

▶ Here, we have independent observations $(\mathbf{x}_i, y_i)$, $i = 1, ..., n$, each follow the (conditional) distribution

$$P(y_i = 1|\mathbf{x}_i) = \frac{1}{1 + \exp(-\beta^T \mathbf{x}_i)} = 1 - P(y_i = 1|\mathbf{x}_i).$$

▶ So, the joint probability function is

$$\prod_{i=1,...,n; y_i=1} p(y_i = 1|\mathbf{x}_i) \prod_{i=1,...,n; y_i=0} p(y_i = 0|\mathbf{x}_i)$$

which can be conveniently written as

$$\prod_{i=1}^{n} \frac{\exp(y_i \beta^T \mathbf{x}_i)}{1 + \exp(\beta^T \mathbf{x}_i)}.$$

The likelihood and log-likelihood

- The likelihood function is the same as the joint probability function, but viewed as a function of β.
- The log-likelihood function is

$$\sum_{i=1}^{n}[y_i\beta^T x_i - \log(1 + \exp(\beta^T \mathbf{x}_i))]$$

- The maximizor is denoted as $\hat{\beta}$, which is the MLE of β based on logistic model.

Example: the default data with all three inputs: balance, income and student

TABLE 4.3. For the Default data, estimated coefficients of the logistic regression model that predicts the probability of default using balance, income, and student status. Student status is encoded as a dummy variable student[Yes], with a value of 1 for a student and a value of 0 for a non-student. In fitting this model, income was measured in thousands of dollars.

	Coefficient	Std.error	t-statistic	p-value
Intercept	-10.8690	0.4923	-22.08	< 0.0001
Balance	0.0057	0.0002	24.74	< 0.0001
Income	0.0030	0.0082	0.37	0.7115
Student[Yes]	-0.6468	0.2362	-2.74	0.0062

The Bayes Theorem

- Suppose there are K, denoted as $1, 2, ..., K$, for the output Y.
- X is the input of p-dimension. Both Y and X are random variables.
- Let $\pi_k = P(Y = k)$.
- Let $f_k(x) = f(x|Y = k)$ be the conditional density function of X given $Y = k$.
- Then, Bayes theorem[1] implies

$$p_k(x) = P(Y = k|X = x) = \frac{\pi_k f_k(x)}{\sum_{j=1}^{K} \pi_j f_j(x)}$$

- We classify a subject with input x into class k, if its $p_k(x)$ is the largest, for $k = 1, ..., K$.

[1]General Bayes theorem: for $A_i \cap A_j = \emptyset, P(\cup_i A_i) = 1$,

$$P(A_k|B) = \frac{P(B|A_k)P(A_k)}{\sum_{j=1}^{K} P(B|A_j)P(A_j)}.$$

Model Assumptions of LDA

X is p-dimensional. $Y = 1, ..., K$, totally K classes. Assume, for $k = 1, ..., K$,

$$X|Y = k \sim N(\mu_k, \Sigma),$$

where μ_p is p-vector and Σ is p by p variance matrix. i.e.,

$$f_k(x) = \frac{1}{(2\pi)^{p/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)\right)$$

Note that we assumed the same Σ for all classes $k = 1, ..., K$.

Computing $p_k(x)$ for LDA

▶ We aim to maximize over k the following

$$p_k(x) = \frac{\pi_k \exp[(-1/2)(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)]}{\sum_{l=1}^{K} \pi_l \exp[(-1/2)(x - \mu_l)^T \Sigma^{-1}(x - \mu_l)]}$$

▶ Comparing $p_k(x)$ is the same as comparing the numerator. Set the k-th score, linear in x,

$$\delta_k(x) = \mu_k^{-1} \Sigma^{-1} x - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k.$$

where the Bayesian classifier is the k with the largest δ_k.

▶ A practical problem: parameter Σ and μ_j and π_j, $j = 1, ..., K$ are usually unknown?

Fisher's Linear Discriminant Analysis

We choose the k-th class such that the following *linear* score function is the largest:

$$\hat{\delta}_k(x) = \hat{\mu}_k^T \hat{\Sigma}^{-1} x - \frac{1}{2}\hat{\mu}_k^T \hat{\Sigma}^{-1}\mu_k + \log \hat{\pi}_k, \qquad (1)$$

where given data $(x_i, y_i), i = 1, ..., n,$

- $\hat{\pi}_k = n_k/n$ is the sample proportion of class k where n_k is the number of subjects in class k
- $\hat{\mu}_k$ is the sample mean of class k

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k} x_i;$$

- $\hat{\Sigma}$ is the pooled (overall) sample covariance

$$\hat{\Sigma} = \frac{1}{n-K} \sum_{k=1}^{K} \sum_{i:y_i=k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T,$$

Example: the default data, a Rare Event problem

- Fit the LDA model to 10,000 training samples gives training error rate of **2.75%**. That is, 275 samples are misclassified.
- Is it good?
- But, actually only **3.33%** of the training samples default, meaning that, a naive classifier that classifies every sample as non-default only has 3.33% error rate.
- So, minimizing misclassification error is misleading here!
- Instead, let's look at the *confusion matrix*.

The Confusion matrix.

TABLE 4.4. A confusion matrix compares the LDA predictions to the true default statuses for the 10, 000 training observations in the Default data set. Elements on the diagonal of the matrix represent individuals whose default statuses were correctly predicted, while off-diagonal elements represent individuals that were misclassified. LDA made incorrect predictions for 23 individuals who did not default and for 252 individuals who did default.

		True default status		
		No	Yes	Total
Predicted	No	9644	252	9896
default status	Yes	23	81	104
	Total	9667	333	10000

Class-specific performance

- Overall misclassification error rate:
  $(252 + 23)/10000 = 2.75\%$.

- Error rate with default people: $252/333 = 75.7\%$, losing a lot!!

- Sensitivity (TPR): $1 - 75.7\% = 24.3\%$

- Error rate within people no-default (FPR): $23/9667 = 0.24\%$.

- Specificity (1-FPR) : $1 - 0.24\% = 99.8\%$

A modification

- The Bayes classifier classifies a subject into class k, if the posterior probability $p_k(x)$ is the largest.

- In this case with two classes (Yes/No), i.e., $K = 2$. Bayes classifier classifies into *default* class if

$$Pr(default = Yes | X = x) > 0.5.$$

- A modification is classifies into *default* class if

$$Pr(default = Yes | X = x) > 0.2.$$

"I would rather kill a thousand by mistake than let one go..."

The classification result of the modification

TABLE 4.5. A confusion matrix compares the LDA predictions to the true default statuses for the 10, 000 training observations in the Default data set, using a modified threshold value that predicts default for any individuals whose posterior default probability exceeds 20 %.

		True default status		
		No	Yes	Total
Predicted	No	9432	138	9570
default status	Yes	235	195	430
	Total	9667	333	10000

Class-specific performance

- Overall error rate: $(235 + 138)/10000 = 3.73\%$. (increased)
- Error rate with default people: $138/333 = 41.4\%$ (lower after modification)
- Sensitivity (TPR): $1 - 41.4\% = $ **58.6%** (Bayes 24.3%)
- Error rate within people no-default (FPR): $235/9667 = 2.43\%$. (increased from 0.24%)
- Specificity: $1 - 2.43\% = 97.57\%$ (Bayes **99.8%**)
- Identification of defaulter (sensitivity) is more important to credit card company!
- This modification may be helpful to the company. A tradeoff of specificity for sensitivity.

The tradeoff

FIGURE 4.7. For the Default data set, error rates are shown as a function of the threshold value for the posterior probability that is used to perform the assignment. The black solid line displays the overall error rate. The blue dashed line represents the fraction of defaulting customers that are incorrectly classified, and the orange dotted line indicates the fraction of errors among the non-defaulting customers.

The ROC curve (Operation characteristic curve)

FIGURE 4.8. A ROC curve for the LDA classifier on the Default data. It traces out two types of error as we vary the threshold value for the posterior probability of default. The actual thresholds are not shown. The true positive rate is the sensitivity: the fraction of defaulters that are correctly identified, using a given threshold value. The false positive rate is 1-specificity: the fraction of non-defaulters that we classify incorrectly as defaulters, using that same threshold value. The ideal ROC curve hugs the top left corner, indicating a high true positive rate and a low false positive rate. The dotted line represents the ROC of random classifier; this is what we would expect if student status and credit card balance are not associated with probability of default. Predicted class

ROC Curve

General confusion matrix

TABLE 4.6. Possible results when applying a classifier or diagnostic test to a population.

		TRUE class		
		- or Null	+ or Non-null	Total
Predicted	- or Null	True Neg. (TN)	False Neg. (FN)	N*
status	+ or Non-null	False Pos. (FP)	True Pos. (TP)	P*
	Total	N	P	

Clarifying the terminology

TABLE 4.7. Important measures for classification and diagnostic testing, derived from quantities in Table 4.6.

Name	Definition	Synonyms
False Pos. rate	FP/N	Type I error, 1 - specificity
True Pos. rate	TP/P	1- Type II error, power, sensitivity, recall
Pos. Pred.value	TP/P*	Precision, 1- false discovery rate
Neg.Pred.value	TN/N*	

- F_1-score is the harmonic mean of precision and recall:

$$F_1 = \frac{2}{\text{precision}^{-1} + \text{recall}^{-1}} = \frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}.$$